

Hybrid Cloud Preprocessing Script

File

`preprocess.py`

Overview

This script preprocesses **Google Cluster Trace–style datasets** into a structured format suitable for **machine learning–based workload classification** in the **Hybrid Cloud Orchestrator** project.

Every time the script runs:

- It randomly samples **100 rows** (for performance and variability).
 - Extracts, normalizes, and derives workload characteristics.
 - Automatically detects and scales **time units**.
 - Estimates missing values, such as **data size**.
 - Saves a **clean, feature-rich CSV** ready for ML training and inference.
-

Input

A CSV file from **Google Cluster Data**, containing columns such as:

```
start_time, end_time, resource_request, assigned_memory,  
priority, scheduling_class, collection_type, collection_name,  
failed, ...
```

These fields typically represent event logs of resource allocations, usage patterns, and job metadata.



Output

A clean dataset with **6 engineered features**:

Column	Description	Type	Example
cpu cores	Number of CPU cores required	Float	0.21, 1.55, 4.22
memory mb	Memory footprint in MB	Float	341.5, 879.0, 2048.0
latency sensitive	Whether the workload is latency-critical	Integer (0/1)	1
execution time	Duration of the task in seconds	Float	3.0, 300.0, 1200.0
data size mb	Estimated data processed in MB	Float	45.2, 912.8
cost traditional	Estimated cost for VM/Kubernetes deployment	Float	0.0005, 0.034
cost serverless	Estimated cost for serverless function deployment	Float	0.00002, 0.067
cost ratio	Ratio (serverless/traditional) — cost efficiency indicator	Float	21.5, 0.9
target platform	ML label: <code>serverless</code> or <code>traditional</code>	String	serverless



How It Works — Step-by-Step

1 Sampling

Randomly selects **100 rows** from the dataset for preprocessing:

```
df_sample = df_raw.sample(n=100)
```

- Ensures quick iteration and balanced workload examples per run.

2 Parsing Resource Requests

Google cluster data encodes CPU and memory as fractional values of the machine's total capacity, e.g.:

```
{'cpus': 0.020660400390625, 'memory': 0.014434814453125}
```

Computation:

CPU cores = `cpu_fraction × machine_cores`

Memory (MB) = `memory_fraction × machine_memory_MB`

Default:

- `machine_cores = 16`
- `machine_mem_mb = 65536 (≈ 64 GB)`

Example:

```
{'cpus': 0.02066, 'memory': 0.01443}
```

becomes:

CPU :
 $0.02066 \times 16 = 0.33$ cores

Memory :
 $0.01443 \times 65536 = 945$ MB

If numeric values are larger than 1, they're treated as **absolute units**.

3 Memory Normalization

If `assigned_memory` exists, it's used to fill missing memory values.

Heuristics:

- Values $\leq 1 \rightarrow$ fractions (\times total machine memory)

- Values < 10000 → MB
 - Very large numbers → bytes (converted via $\div 1024^2$)
-

4 Execution Time Calculation

Google trace timestamps (`start_time`, `end_time`) are large because they're often in **nanoseconds**.

The script automatically detects the unit and normalizes it to **seconds**.

Detection Logic:

```
if median_diff > 1e9: # nanoseconds
    diff = diff / 1e9

elif median_diff > 1e6: # microseconds
    diff = diff / 1e6

elif median_diff > 1e3: # milliseconds
    diff = diff / 1e3
```

Mathematical Formula:

Execution Time (s) = $(\text{end_time} - \text{start_time}) / k$

Where $k \in \{1, 10^3, 10^6, 10^9\}$ depending on detected units.

Example:

```
start_time = 2.748E+11
end_time   = 2.751E+11
diff = 3E+08 → divided by 1E+9 = 3.0 seconds
```

5 Latency Sensitivity Detection

Heuristic rules based on job metadata:

Field	Rule	Meaning
priority	$\leq 2 \rightarrow 1$	High priority → latency-sensitive
scheduling_class	$\leq 1 \rightarrow 1$	Real-time or interactive
collection_name / collection_type	Contains “api”, “realtime”, “interactive”, “latency”, “foreground” → 1	Textual cues for latency-bound tasks
Otherwise	0	Background or batch jobs

Formula:

Formula:

$$\text{latency_sensitive} = \begin{cases} 1, & \text{if priority} \leq 2 \text{ or scheduling_class} \leq 1 \text{ or contains keyword} \\ 0, & \text{otherwise} \end{cases}$$

6 Data Size Estimation

Google dataset doesn't contain explicit data size..

Formula:

$$\text{data_size_mb} = \frac{\text{memory_mb} \times \text{execution_time}}{100}$$

This keeps the scale realistic and consistent (reduces very large magnitudes).

Example:

$$\text{memory_mb} = 1024, \text{execution_time} = 5$$

$$\text{data_size_mb} = \frac{1024 \times 5}{100} = 51.2$$

7 💰 Cost Estimation

This section estimates both **Traditional** and **Serverless** costs for each workload.

■ Traditional Compute (VM / Kubernetes)

Charged per core-hour and GB-hour:

$$\text{Cost}_{\text{traditional}} = (\text{CPU cores} \times P_{\text{CPU}} + \text{Memory (GB)} \times P_{\text{MEM}}) \times \frac{\text{Execution Time (s)}}{3600}$$

where:

- $P_{\text{CPU}} = 0.0316$ \$/core-hour
- $P_{\text{MEM}} = 0.0045$ \$/GB-hour

■ Serverless Compute (Lambda / Cloud Functions)

Charged per GB-second + per request:

$$\text{Cost}_{\text{serverless}} = P_{\text{req}} + P_{\text{GB-sec}} \times \text{Memory (GB)} \times \text{Execution Time (s)}$$

where:

- $P_{\text{GB-sec}} = 0.00001667$
- $P_{\text{req}} = 0.0000002$

Pricing is not random.

Constant	Description	Value (USD)	Derived From
P_CPU	Cost per vCPU-hour	0.0316	GCP/AWS EC2 averages
P_MEM	Cost per GB-hour	0.0045	GCP/AWS VM averages
P_GB_SEC	Serverless compute price	0.00001667	AWS Lambda standard rate
P_REQ	Request charge (per invocation)	0.0000002	AWS Lambda request pricing

⚖️ Cost Ratio

A cost comparison metric:

A cost comparison metric:

$$\text{cost_ratio} = \frac{\text{cost_serverless}}{\text{cost_traditional} + 10^{-9}}$$

Ratio Value	Interpretation	Meaning in Context
< 1	Serverless cheaper	The same job costs less on serverless than on a VM. → Ideal for bursty, short workloads.
= 1	Equal cost	Both environments cost roughly the same. Decision may depend on latency or management preference.
> 1	Traditional cheaper	Serverless costs more than running on a VM — typically long or compute-heavy jobs.

8 Failure Normalization

Converts textual failure states into boolean flags.

Input	Output
1, true, y, fail, failed, failure	True
0, empty, others	False

This is used in target classification.

9 Target Platform Classification

The **core ML label** (`serverless` or `traditional`) is determined heuristically:

Logic:

A workload is considered **serverless** if:

- Short-lived
- Lightweight
- Memory-efficient

Mathematical Rules:

$$\text{serverless if: } \begin{cases} \text{execution_time} \leq 300 \text{ seconds} \\ \text{cpu_cores} \leq 2 \\ \text{memory_mb} \leq 2048 \\ \text{failed} = \text{False} \end{cases}$$

Otherwise:

target platform = "traditional"

Sampling and Saving

- Randomly samples 100 rows each run (new subset each execution).

Writes processed dataset to:

`workload_dataset_sample.csv`

- File automatically **overwrites** if it already exists (for versioning simplicity).
-



Example Output

cpu cores	memory mb	latency sensitivity	execution time	data size mb	cost traditional	cost serverless	cost ratio	target platform
0.5	512	1	5	25.6	0.000002	0.000043	21.5	serverless
2	2048	0	300	6144	0.0016	0.102	63.75	traditional
8	8192	0	3600	294912	0.236	491.52	2081.0	traditional



Default Parameters

Parameter	Default	Description
--cores	16	Number of CPU cores per machine
--mem_mb	65536	Machine total memory in MB
--input	<i>required</i>	Input dataset CSV
--output	workload_dataset_sample100.csv	Output file
--seed	None	Random seed for reproducible sampling

Example Command

```
python preprocess.py --input Sample_Dataset.csv --output  
workload_dataset_sample.csv
```

or, with explicit configuration:

```
python sample_and_preprocess.py --input google_dataset.csv --output  
workload_dataset_sample100.csv --cores 32 --mem_mb 131072
```

Summary of All Math & Logic

Feature	Formula
CPU cores	<code>cpu_fraction × machine_cores</code>
Memory (MB)	<code>mem_fraction × machine_memory_MB</code>
Execution time (s)	$(\text{end_time} - \text{start_time})/k$, where $k = 1, 10^3, 10^6, 10^9$
Latency sensitive	1 if <code>priority ≤ 2</code> or <code>scheduling_class ≤ 1</code>
Data size (MB)	<code>memory_mb × execution_time/100</code>
Cost traditional	$(\text{CPU} \times P_{CPU} + \text{MEM(GB)} \times P_{MEM}) \times \text{runtime(hr)}$
Cost serverless	$P_{req} + P_{GB-sec} \times \text{MEM(GB)} \times \text{time(s)}$
Cost ratio	<code>cost_serverless/cost_traditional</code>
Target platform	Serverless if time ≤ 300s, CPU ≤ 2, Mem ≤ 2GB, else Traditional

Purpose in the Project

This script is the **first stage of the intelligent hybrid cloud orchestration pipeline** — it converts raw Google cluster traces into **structured, ML-ready workload features**, enabling the **ML classifier** to learn how to route workloads between:

-  **Serverless platforms** (e.g., AWS Lambda, GCP Functions)
-  **Traditional infrastructure** (e.g., EC2, Kubernetes)